




## PREDICCIÓN DE VIRALIDAD DE TWEETS

Marc Pascual  
Alejandro Lobo



# Introducción

Fuente del dataset y  
objetivo



¿Porque lo hemos escogido?





# Introducción

## Análisis general


```
Python  
data.isnull().any()
```

```
Python ▾  
Date False  
Time False  
Tweet_Text False  
Type False  
Media_Type True  
Hashtags True  
Tweet_Id False  
Tweet_Url False  
twt_favourites_IS_THIS_LIKE_QUESTION_MARK False  
Retweets False  
Unnamed: 10 True  
Unnamed: 11 True  
dtype: bool
```



# Introducción

## Relación entre las elecciones y los likes y RT



- Mas likes y RT cerca de las elecciones
- Decidimos ignorar este hecho



# Preprocesamiento

## Importación del dataset

### CARGAR LOS DATOS

```
pandas.set_option('display.max_columns', None)
pandas.set_option('display.expand_frame_repr', False)
pandas.set_option('display.precision', 3)
df = pandas.read_csv('data.csv', sep=',', na_values="")
print(df.head())
```

	Date	Time	Tweet_Text	Type	Media_Type	Hashtags	Tweet_Id	Tweet_Url	twt_favourites_IS_THIS_QUESTION_MARK	Retweets	Unnamed: 10	Unnamed: 11	
0	16-11-11	15:26:37	Today we express our deepest gratitude to all ...	text	photo	ThankAVet	7.970e+17	<a href="https://twitter.comrealDonaldTrump/status/797...">https://twitter.com/realDonaldTrump/status/797...</a>		127213	41112	NaN	NaN
1	16-11-11	13:33:35	Busy day planned in New York. Will soon be mak...	text	NaN	NaN	7.970e+17	<a href="https://twitter.comrealDonaldTrump/status/797...">https://twitter.com/realDonaldTrump/status/797...</a>		141527	28654	NaN	NaN
2	16-11-11	11:14:20	Love the fact that the small groups of protest...	text	NaN	NaN	7.970e+17	<a href="https://twitter.comrealDonaldTrump/status/797...">https://twitter.com/realDonaldTrump/status/797...</a>		183729	50039	NaN	NaN
3	16-11-11	2:19:44	Just had a very open and successful presidenti...	text	NaN	NaN	7.970e+17	<a href="https://twitter.comrealDonaldTrump/status/796...">https://twitter.com/realDonaldTrump/status/796...</a>		214001	67010	NaN	NaN
4	16-11-11	2:10:46	A fantastic day in D.C. Met with President Oba...	text	NaN	NaN	7.970e+17	<a href="https://twitter.comrealDonaldTrump/status/796...">https://twitter.com/realDonaldTrump/status/796...</a>		178499	36688	NaN	NaN



# Preprocesamiento

## Limpieza de tweets

### LIMPIAR EL TEXTO

```
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import SnowballStemmer
# Preparación de las herramientas de preprocesamiento de texto
nltk.download('stopwords')
stop = set(stopwords.words('english'))
sno = SnowballStemmer('english')

def cleanhtml(sentence):
    cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext

def cleanpunc(sentence):
    cleaned = re.sub(r'[!|!|\`|"|"|#]',r'',sentence)
    cleaned = re.sub(r'[.,!|,!]|(|\|/|)',r' ',cleaned)
    return cleaned

# Procesamiento del texto de los tweets
final_string = []
for sent in df['Tweet_Text'].values:
    filtered_sentence = []
    sent = cleanhtml(sent)
    for w in sent.split():
        cleaned_words = cleanpunc(w)
        if cleaned_words.isalpha() and len(cleaned_words) > 2:
            if cleaned_words.lower() not in stop:
                stemmed_word = sno.stem(cleaned_words.lower())
                filtered_sentence.append(stemmed_word)
    final_string.append(" ".join(filtered_sentence))

# Añadiendo la columna de tweets limpios al DataFrame
df['cleaned_tweet'] = final_string

# Visualización de los primeros registros del DataFrame modificado
print(df['cleaned_tweet'].head())
print(df['cleaned_tweet'].iloc[0])
```

### Antes

... Today we express our deepest gratitude to all those who have served in our armed forces. #ThankAVet <https://t.co/wPk70WpK8Z>

### Después

today express deepest gratitud serv arm thankavet



# Preprocesamiento

## Creación de Viralidad

### CREAR PORCENTAJE DE MUCHO Y POCO

```
media_likes= df["twt_favourites_IS_THIS_LIKE_QUESTION_MARK"].median()
media_retweets=df["Retweets"].median()
df['Viral'] = df.apply(lambda x: 'mucho' if (x['twt_favourites_IS_THIS_LIKE_QUESTION_MARK'] + x['Retweets']) >= (media_likes + media_retweets) * 1.05
| else ('poco'), axis=1)
print(df["Viral"].value_counts())
```

[13]

```
...
Viral
poco    3773
mucho   3602
Name: count, dtype: int64
```



# Preprocesamiento

## Creación de Viralidad

### CREAR PORCENTAJE DE MUCHO Y POCO

```
media_likes= df["twt_favourites_IS_THIS_LIKE_QUESTION_MARK"].median()
media_retweets=df["Retweets"].median()
df['Viral'] = df.apply(lambda x: 'mucho' if (x['twt_favourites_IS_THIS_LIKE_QUESTION_MARK'] + x['Retweets']) >= (media_likes + media_retweets) * 1.05
| else ('poco'), axis=1)
print(df["Viral"].value_counts())
```

[13]

```
...
Viral
poco    3773
mucho   3602
Name: count, dtype: int64
```



# Preprocesamiento

## Eliminación de columnas

### ELIMINAR COLUMNAS INESESCAREAS

```
df.drop('Tweet_Text', axis=1, inplace=True)
df.drop('Tweet_Id', axis=1, inplace=True)
df.drop('Tweet_Url', axis=1, inplace=True)
df.drop('Date', axis=1, inplace=True)
df.drop('Time', axis=1, inplace=True)
df.drop('Media_Type', axis=1, inplace=True)
df.drop('Type', axis=1, inplace=True)
df.drop('Hashtags', axis=1, inplace=True)
df.drop('twt_favourites_IS_THIS_QUESTION_MARK', axis=1, inplace=True)
df.drop('Retweets', axis=1, inplace=True)
df.drop('Unnamed: 10', axis=1, inplace=True)
df.drop('Unnamed: 11', axis=1, inplace=True)
y=df["Viral"].values
# Definir X como todas las columnas excepto 'Viral'
X= df.drop('Viral', axis=1).values
print(df[0:5])
```

## Dataset Limpio

	cleaned_tweet	Viral
0	today express deepest gratitud serv arm thankavet	mucho
1	busi day plan new soon make import decis peopl...	mucho
2	love fact small group protest last night passi...	mucho
3	open success presidenti profession incit unfair	mucho
4	fantast day met presid obama first realli good...	mucho



# Preprocesamiento

## Creación de columnas

GENERAR COLUMNAS POR PALABRAS I FILTRAR POR UMBRAL

```
from sklearn.feature_extraction.text import CountVectorizer

count_vectorizer = CountVectorizer()
X = count_vectorizer.fit_transform( X[:, -1])
feature_names = count_vectorizer.get_feature_names_out()
X = pandas.DataFrame(X.toarray(), columns=feature_names)
word_frequencies = X.sum(axis=0)
filtered_words = word_frequencies[word_frequencies >= 5]
X = X[filtered_words.index]
X.to_csv('X.csv', index=False)
df.drop('cleaned_tweet', axis=1, inplace=True)
df.to_csv('y.csv', index=False)
print(X[0:5])
print(y[0:5])
print(X.shape)
```

## Dataset X e y

	abc	abl	abolish	absolut	accept	accord	act
0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0
2	0	0	0	0	0	0	0
3	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0

['mucho' 'mucho' 'mucho' 'mucho' 'mucho']  
(7375, 1477)



# Criterios de evaluación

## Evaluar División Fija

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Ratios de Training, Validación y Test
split_ratios = [(0.7, 0.1), (0.6, 0.1), (0.5, 0.2)] # Train, Validation ratios
results = []

for train_ratio, val_ratio in split_ratios:
    # Primer paso: dividir en entrenamiento y temporal
    X_train, X_temp, y_train, y_temp = train_test_split(X, y, test_size=1-train_ratio, random_state=42)

    # Segundo paso: dividir el conjunto temporal en validación y prueba
    X_val, X_test, y_val, y_test = train_test_split(
        X_temp, y_temp, test_size=val_ratio/(val_ratio + (1-train_ratio-val_ratio)), random_state=42
    )

    # Crear y entrenar el modelo
    model = DecisionTreeClassifier()
    model.fit(X_train, y_train)

    # Evaluar el modelo
    val_accuracy = accuracy_score(y_val, model.predict(X_val))
    test_accuracy = accuracy_score(y_test, model.predict(X_test))

    # Guardar resultados
    results.append({
        'train_ratio': train_ratio,
        'val_ratio': val_ratio,
        'test_ratio': 1-train_ratio-val_ratio,
        'val_accuracy': val_accuracy,
        'test_accuracy': test_accuracy
    })
# Imprimir resultados
for res in results:
    print(f"Train: {res['train_ratio']*100}%, Validation: {res['val_ratio']*100}%, Test: {res['test_ratio']*100}%")
    print(f"Validation Accuracy: {res['val_accuracy']:.4f}, Test Accuracy: {res['test_accuracy']:.4f}")
    print()
```

## Valores

Train: 70.0%, Validation: 10.0%, Test: 20.0%  
Validation Accuracy: 0.6807, Test Accuracy: 0.6707

Train: 60.0%, Validation: 10.0%, Test: 30.0%  
Validation Accuracy: 0.6786, Test Accuracy: 0.6504

Train: 50.0%, Validation: 20.0%, Test: 30.0%  
Validation Accuracy: 0.6655, Test Accuracy: 0.6551



# Criterios de evaluación

## División de los datos

```
Python
(X_train, X_test, y_train, y_test) = train_test_split(X, y2, test_size=.3,
random_state=1)
```



# Criterios de evaluación

## Single vs K fold cross validation

```
%matplotlib inline
import pandas
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import sklearn
import sklearn.datasets as ds
import sklearn.model_selection as cv
import sklearn.neighbors as nb
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score
from sklearn.metrics import f1_score, classification_report, confusion_matrix
from sklearn.feature_selection import SelectKBest, f_classif
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

# Cargar los datos (esto ya lo tienes)
X = pandas.read_csv('X.csv', sep=',', na_values="")
y = pandas.read_csv('y.csv', sep=',', na_values="")

# 1. **Single Fold Validation** - División entre datos de entrenamiento y test (70-30)
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3, random_state=42)

# Inicializar el clasificador KNN
knn = KNeighborsClassifier(n_neighbors=5)

# Entrenamiento del modelo
knn.fit(X_train, y_train)

# Predicciones
y_pred = knn.predict(X_val)

# Calcular la precisión y otras métricas
accuracy = accuracy_score(y_val, y_pred)
print("\nSingle Fold Validation Results:")
print(f"Accuracy: {accuracy}")

# 2. **K-Fold Cross-Validation** - Usando 5 pliegues
cv = 10

# Inicializar KNN nuevamente
knn = KNeighborsClassifier(n_neighbors=5)

# Realizar K-Fold Cross-Validation
scores = cross_val_score(knn, X, y, cv=cv, scoring='accuracy')

# Mostrar los resultados de la validación cruzada
print("\nK-Fold Cross-Validation Results:")
print(f"Average accuracy: {scores.mean()}")
```

## Valores

Single Fold Validation Results:  
Accuracy: 0.6037053773158608



K-Fold Cross-Validation Results:  
Average accuracy: 0.5919848282607657



# Criterios de evaluación

## Métricas de evaluación

```
Python
print(sklearn.metrics.accuracy_score(y_test, pred))
print(classification_report(y_test, pred))
cm = (confusion_matrix(y_test, pred))
plt.figure(figsize=(6, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["Poco",
"Mucho"], yticklabels=["Poco", "Mucho"])
plt.show()
# Reporte de clasificación
report = sklearn.metrics.classification_report(y_test, pred,
output_dict=True)
# Visualización de métricas globales del reporte
metrics = ['precision', 'recall', 'f1-score']
report_df = pandas.DataFrame(report).T
report_df = report_df[metrics]
# Plot de métricas
report_df.iloc[:-3].plot(kind='bar', figsize=(10, 6), colormap='viridis',
edgecolor='black')
plt.title("Precisión, Recall y F1-score por clase")
plt.ylabel("Score")
plt.xlabel("Clases")
plt.xticks(rotation=0)
plt.grid(axis='y')
plt.show()
```





# Naive Bayes

## Independencia de Variables en Naïve Bayes

Python

```
from sklearn.feature_selection import mutual_info_classif
from sklearn.metrics import pairwise_distances
word_correlations = np.corrcoef(X, rowvar=False)
top_correlated_pairs = []
threshold = 0.5
for i in range(word_correlations.shape[0]):
    for j in range(i + 1, word_correlations.shape[1]):
        if abs(word_correlations[i, j]) > threshold:
            top_correlated_pairs.append((i, j, word_correlations[i, j]))
print("Top Correlated Pairs")
print(len(top_correlated_pairs))
```

Python

```
Top Correlated Pairs
53
```

## Tenemos suficientes datos

```
print(X.shape)
print(y.shape)
```

[2]

```
... (7375, 1477)
(7375, 1)
```

poco 3773  
mucho 3602



# Naive Bayes

## Decisión de utilizar MultinomialNB

The infographic illustrates three types of Naive Bayes models:

- Bernoulli Naive Bayes**: For binary or boolean features. Example table:

	HAS 4 LEGS?	CAN FLY?	LAY EGGS?
1	0	1	
0	1	1	
0	1	1	
0	1	0	
0	0	1	

- Multinomial Naive Bayes**: For discrete features (like word count). Example table:

WORD	the	but	is
4	1	2	
3	0	0	
3	0	0	
1	0	4	
0	2	3	


- Gaussian Naive Bayes**: For continuous, real-valued attributes. Example table:

AGE	WEIGHT	HEIGHT
17.4	56.5	145.2
25.4	71.2	170.4
18.8	70.3	164.5
21.1	51.2	140.5
20.9	81.5	182.2



# Naive Bayes

## Optimización del umbral



- Sensible a partición de datos
- umbral menor a 0,5



# Naive Bayes

## Resultados

Con umbral

Sin umbral

Python

	precision	recall	f1-score	support
0.0(mucho)	0.50	0.63	0.66	1066
1.0(poco)	0.66	0.92	0.77	1147
accuracy			0.72	2213
macro avg	0.76	0.71	0.70	2213
weighted avg	0.75	0.72	0.70	2213

Python

0.7144148215092635	precision	recall	f1-score	support
0.0	0.72	0.67	0.69	1066
1.0	0.71	0.76	0.73	1147
accuracy			0.71	2213
macro avg	0.71	0.71	0.71	2213
weighted avg	0.71	0.71	0.71	2213




# KNN

## Weighting" o "No weighting

```
lr = []
for ki in range(1,30,2):
    cv_scores = cross_val_score(nb.KNeighborsClassifier(n_neighbors=ki), X=X_train, y=y_train, cv=10)
    lr.append(np.mean(cv_scores))
plt.plot(range(1,30,2),lr,'b',label='No weighting')

lr = []
for ki in range(1,30,2):
    cv_scores = cross_val_score(nb.KNeighborsClassifier(n_neighbors=ki,weights='distance'), X=X_train, y=y_train, cv=10)
    lr.append(np.mean(cv_scores))
plt.plot(range(1,30,2),lr,'r',label='Weighting')
plt.xlabel('k')
plt.ylabel('Accuracy')
plt.legend(loc='upper right')
plt.grid()
plt.tight_layout()

plt.show()
```





# KNN

## Mejor parámetro k

```
from sklearn.model_selection import GridSearchCV
params = {'n_neighbors':list(range(1,30,2)), 'weights':('distance','uniform')}
knc = nb.KNeighborsClassifier()
clf = GridSearchCV(knc, param_grid=params, cv=10, n_jobs=-1) # If cv is integer, by default is Stratified
clf.fit(X_train, y_train)
print("Best Params=", clf.best_params_, "Accuracy=", clf.best_score_)
```

---

Best Params= {'n\_neighbors': 11, 'weights': 'distance'} Accuracy= 0.6859681675738083




# KNN

## Resultados

Python  
precision recall f1-score support

	mucho	poco	precision	recall
precision	0.74	0.65	0.52	0.61
recall			0.83	0.73
f1-score	0.68	0.67	0.61	0.61
support	1866	1147	2213	2213
accuracy			0.68	0.68
macro avg	0.70	0.68	0.67	0.67
weighted avg	0.69	0.68	0.67	0.67





# Decision Trees

## Evaluación del Modelo

```
from sklearn import tree

clf = tree.DecisionTreeClassifier(criterion='entropy', min_impurity_decrease=0.0001)
pred = clf.fit(X_train, y_train).predict(X_test)
print("Accuracy on test set:", sklearn.metrics.accuracy_score(y_test, pred))
epsilon = sklearn.metrics.accuracy_score(y_test, pred)
print(sklearn.metrics.classification_report(y_test, pred))
conf_matrix = sklearn.metrics.confusion_matrix(y_test, pred)
plt.figure(figsize=(8, 6))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=np.unique(y), yticklabels=np.unique(y))
plt.show()

# Reporte de clasificación
report = sklearn.metrics.classification_report(y_test, pred, output_dict=True)
# Visualización de métricas globales del reporte
metrics = ['precision', 'recall', 'f1-score']
report_df = pandas.DataFrame(report).T
report_df = report_df[metrics]
# Plot de métricas
report_df.iloc[:3].plot(kind='bar', figsize=(10, 6), colormap='viridis', edgecolor='black')
plt.title("Precisión, Recall y F1-score por clase")
plt.ylabel("Score")
plt.xlabel("Clases")
plt.xticks(rotation=0)
plt.grid(axis='y')
plt.show()


print("Confidence interval: ", proportion_confint(count=epsilon*X_test.shape[0], nobs=X_test.shape[0], alpha=0.05, method='binom_test'))
```




# Decision Trees

## min\_impurity\_decrease

Impurity = 0.001




Impurity = 0.0001





# Decision Trees


## Interpretación del Arbol





# Decision Trees


## Interpretación del Arbol





# Decision Trees



## Interpretación del Arbol





# Decision Trees

## Resultados






# Support Vector Machines

## Elección del Kernel y parámetros

Kernel lineal



Kernel Polinomial y RBF





# Support Vector Machines

## Método para acelerar el training

Reducción número de  
características

```
print(X.shape)
print(y.shape)
```

---

```
(295, 34)
(295, 1)
```

Reducción de parámetros

```
# Lista de valores de C a probar. Usualmente
Cs = np.logspace(-3, 5, num=3, base=10.0)
```



# Support Vector Machines

## Resultados

### Kernel Lineal

- Mejor parámetro C encontrado: 10.0
- Exactitud en el conjunto de test: 64.04%
- Número de soportes: **152** (73.78% de los datos de entrenamiento)
- Intervalo de confianza en validación cruzada para el mejor C:
- Precisión promedio: 58.28%
- Intervalo de confianza (95%): [50.86%, 65.71%]

### Kernel Polinomial

- Mejor combinación de parámetros: C = 10000.0, Grado = 3
- Exactitud en el conjunto de test: 58.43%
- Número de soportes: **137** (66.50% de los datos de entrenamiento)
- Intervalo de confianza en validación cruzada para el mejor C:
- Precisión promedio: 55.34%
- Intervalo de confianza (95%): [54.49%, 56.19%]

### Kernel RBF (Radial Basis Function)

- Mejor combinación de parámetros: C = 10000.0, Gamma = 1.0
- Exactitud en el conjunto de test: 61.80%
- Número de soportes: **178** (86.41% de los datos de entrenamiento)
- Intervalo de confianza en validación cruzada para el mejor C y Gamma:
- Precisión promedio: 52.92%
- Intervalo de confianza (95%): [52.05%, 53.78%]



# Meta Methods

## Majority Voting



Hard



Soft





# Meta Methods

## Resultados

HARD

```
import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)

eclf = VotingClassifier(estimators=[('nb', clf1), ('knn3', clf2), ('dt', clf3)], voting='hard')
scores = cross_val_score(eclf, X, y, cv=cv, scoring='accuracy')
print("Accuracy: %0.3f [%s]" % (scores.mean(), "Majority Voting"))
```

1

Accuracy: 0.636 [Majority Voting]

SOFT

```
eclf = VotingClassifier(estimators=[('nb', clf1), ('knn3', clf2), ('dt', clf3)], voting='soft', weights=[2.5,1.5,2])
scores = cross_val_score(eclf, X, y, cv=cv, scoring='accuracy')
print("Accuracy: %0.3f [%s]" % (scores.mean(), "Weighted Voting"))
```


1

Accuracy: 0.635 [Weighted Voting]



# Meta Methods


Bagging, Boosting, Random Forest





# Meta Methods

## Resultados





# Conclusiones

## Cross Validation

Cross Validation	Accuracy
Single-Fold CV	60.37%
K-Fold CV	59.20%



# Conclusiones

## Learning Methods

Model/Classificador	Accuracy
Naïve Bayes	72%
KNN	68%
Decision Tree	67%
SVM	58%
Meta-learning (Random Forest)	71%

Model/Classificador	Class	precision	recall	f1-score
Naïve Bayes	Mucho	0.72	0.67	0.69
	Poco	0.71	0.76	0.73
KNN	Mucho	0.74	0.52	0.61
	Poco	0.65	0.83	0.73
Decision Tree	Mucho	0.67	0.68	0.68
	Poco	0.68	0.68	0.68
SVM	Mucho	0.68	0.40	0.51
	Poco	0.54	0.79	0.64
Meta-learning (Random Forest)	Mucho	0.72	0.70	0.71
	Poco	0.71	0.73	0.72



# Conclusiones

## Intervalos de confianza

Model/Classificador	Intervalos de Confianza
Naïve Bayes	(0.696, 0.733)
KNN	(0.662, 0.701)
Decision Tree	(0.657, 0.696)
SVM	(0.509, 0.657)
Meta-learning (Random Forest)	(0.636, 0.658) <input type="checkbox"/>